

BEST-OF-N CAN AMPLIFY SPARSE-VIEW GEOMETRY ERRORS IN 3D SCENE WORLD MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

Learned 3D scene world models increasingly represent futures as neural fields, Gaussian splats, point clouds, or occupancy volumes. A common inference pattern is to sample multiple candidate futures and select the candidate with the best proxy score. This paper studies a narrow failure mode of that pattern: when the proxy observes only sparse views, Best-of-N selection can prefer candidates that match the scored views while becoming less consistent in hidden 3D regions. We build a deterministic synthetic voxel benchmark that isolates the mechanism. Naive sparse-view Best-of-N improves proxy score from 0.791 at $N = 1$ to 1.000 at $N = 64$, but true 3D IoU falls from 0.520 to 0.340 and hidden-region error rises from 0.442 to 0.929. A simple reranker using ensemble geometry consensus, visual-hull volume plausibility, and surface-complexity penalties recovers much of the hidden geometry, reaching 0.888 true IoU at $N = 16$. The result is not a claim that all neural scene models fail this way; it is a runnable diagnostic for an under-tested interaction between sparse projective scoring and multi-sample inference.

1 INTRODUCTION

Best-of-N inference is attractive for world models: sample several plausible future scenes, score them with a task proxy, and execute or report the best. In 3D scene representations this proxy often comes from sparse observations: rendered views, depth maps, partial point clouds, or task-relevant camera rays. The proxy is cheaper and more available than full 3D ground truth, but it is not injective. Many different 3D scenes can share the same silhouettes and visible front surfaces.

This paper asks what happens when that ambiguity meets Best-of-N selection. The answer in our controlled benchmark is uncomfortable: increasing N can make the selected candidate look better under the sparse proxy while making it worse as a 3D scene. The selector finds rare view impostors. These candidates are excellent under the scored views and poor in hidden regions.

Our contribution is a mechanism-specific diagnostic rather than a new renderer:

- a synthetic 3D scene-world benchmark with primitive occupancy scenes, stochastic candidate futures, sparse orthographic proxy views, and hidden geometry labels for evaluation;
- diagnostics for proxy score, true 3D IoU, hidden-region error, selected-candidate diversity, view-impostor rate, and the geometric exploitation gap;
- a construction-level proposition showing why sparse-view equivalence enables N -dependent amplification;
- a simple coverage-aware reranker that reduces the exploitation gap in the benchmark without accessing hidden ground truth.

2 RELATED WORK

Neural scene representations. NeRF Mildenhall et al. (2020), Scene Representation Networks Sitzmann et al. (2019), sparse-view NeRF methods such as pixelNeRF Yu et al. (2021), IBRNet Wang et al. (2021), RegNeRF Niemeyer et al. (2022), and SparseNeRF Wang et al. (2023) study

how to represent or infer scenes from limited views. Gaussian Splatting Kerbl et al. (2023) and follow-up geometry-aware splatting methods such as SuGaR Guedon and Lepetit (2024) and 2DGS Huang et al. (2024) provide explicit, fast renderable scene representations. These papers motivate the architecture class, but they do not isolate Best-of-N selection over candidate scene futures.

Occupancy and scene world models. Occupancy Networks Mescheder et al. (2019), Convolutional Occupancy Networks Peng et al. (2020), MonoScene Cao and de Charette (2022), VoxFormer Li et al. (2023), and OccWorld Zheng et al. (2024) address 3D completion and occupancy prediction. Our benchmark uses voxels because they make hidden-region error directly measurable; the question is not how to build the best occupancy predictor, but how a sparse proxy behaves when choosing among multiple predicted occupancies.

Best-of-N and proxy overoptimization. Multi-hypothesis prediction and best-of-many objectives Guzman-Rivera et al. (2012); Bhattacharyya et al. (2018) encourage diverse outputs. Reward and verifier literature shows that selecting against an imperfect proxy can overoptimize that proxy Cobbe et al. (2021); Lightman et al. (2023). Our work specializes this concern to 3D projective ambiguity: sparse views create equivalence classes of hidden geometry.

Classical multiview geometry. Visual hull and space-carving results Laurentini (1994); Kutulakos and Seitz (2000) already show that sparse views underdetermine shape. We use this as the mechanism, not as a novelty claim. The new empirical question is how often Best-of-N selection finds such ambiguous impostors as N grows.

3 BENCHMARK

Scenes are binary occupancy grids composed of random boxes and ellipsoids. A target scene is treated as the future scene to be predicted. A simulated world-model sampler emits N candidate scenes from three modes:

- **faithful:** small shifts, deletions, and local additions around the target;
- **drift:** larger global shifts and noise;
- **view impostor:** copies first visible surfaces under the sparse proxy views and randomizes hidden voxels behind those surfaces.

The sparse proxy observes two orthographic axes. For each view it scores silhouette IoU and first-hit depth agreement:

$$s_{\text{proxy}}(x, y) = \frac{1}{|V|} \sum_{v \in V} 0.7 \text{IoU}(\text{sil}_v(x), \text{sil}_v(y)) + 0.3 \text{depthsim}_v(x, y).$$

Evaluation uses true 3D IoU and hidden-region IoU, where the hidden region is the target visual hull excluding voxels visible from the proxy views.

4 MECHANISM

Proposition 1 (Sparse-view Best-of-N amplification). *For a finite sparse set of projection views, there exist distinct 3D occupancies with identical proxy score and different true 3D IoU. If a sampler independently produces a proxy-dominating hidden-inconsistent candidate with probability $p > 0$, then the probability that naive Best-of-N sees at least one such candidate is $1 - (1 - p)^N$.*

Sketch. Copy the first occupied voxel on every scored ray from a target scene. Any occupancy added or removed strictly behind those copied first hits does not change the scored silhouettes or first-hit depths. Therefore sparse proxy score cannot distinguish the target from these hidden variants. Under independent sampling, the probability of at least one hidden-inconsistent variant among N samples is the complement of sampling none. If that variant has higher proxy score than faithful alternatives, the selector chooses it.

The proposition is deliberately narrow. It assumes binary occupancy, orthographic first-hit views, and independent samples. Its purpose is to justify the diagnostic tested below, not to prove a universal theorem for all renderers.




Figure 1: Naive Best-of-N improves the sparse proxy while reducing true 3D consistency. The coverage-aware reranker trades a small amount of proxy score for much better 3D IoU and lower exploitation gap.

Table 1: Full synthetic run with 72 scenes. Values are means over scenes.

method	N	proxy	true IoU	hidden error	gap
naive	1	0.791	0.520	0.442	0.270
naive	8	0.995	0.400	0.841	0.594
naive	16	0.999	0.350	0.915	0.650
naive	64	1.000	0.340	0.929	0.660
repaired	1	0.791	0.520	0.442	0.270
repaired	8	0.953	0.814	0.200	0.138
repaired	16	0.962	0.888	0.121	0.074
repaired	64	0.978	0.851	0.183	0.127

5 REPAIR

The repair keeps the candidate set fixed and changes only the selection score:

$$s_{\text{repair}} = 0.58s_{\text{proxy}} + 0.27s_{\text{consensus}} + 0.22s_{\text{volume}} - 0.12s_{\text{surface}}.$$

The consensus term is IoU with a thresholded ensemble occupancy. The volume term penalizes candidates whose occupied mass is implausible under the visual hull induced by the sparse views. The surface term penalizes thin shell-like occupancy. No term uses hidden ground-truth voxels.

6 RESULTS

The headline result is the divergence between proxy and true 3D consistency. For naive Best-of-N, proxy score reaches 1.000 at $N = 32$ and $N = 64$, while true IoU falls to 0.340. Hidden-region error rises to 0.929. The selected candidate is a view impostor in 100% of scenes at $N = 32$ and $N = 64$.

The repair is not perfect, but it targets the mechanism. It reaches 0.888 true IoU at $N = 16$ and keeps the exploitation gap at 0.074. At $N = 64$, the repair still selects some impostors, but true IoU remains 0.851, far above the naive 0.340.

7 LIMITATIONS

This is a synthetic mechanism study. The candidate sampler is simulated, the proxy is a simplified silhouette/depth score, and the repair is hand-specified. The benchmark does not establish failure rates for real NeRF, Gaussian splatting, point-cloud, or autonomous-driving occupancy systems. The next submission-quality step is to attach the diagnostic to a real neural-scene generator and evaluate with held-out 3D reconstructions or multi-view captures.




Figure 2: Naive selection increases hidden-region error and collapses top-proxy diversity as N grows.




Figure 3: Example projections. The naive selected candidate matches sparse scored views but differs in unscored geometry; the repaired candidate preserves more global occupancy.

8 REPRODUCIBILITY

The repository contains deterministic seeds, tests, smoke and full presets, CSV tables, and generated figures. The main commands are:

```
python -m pytest
python experiments/run_synthetic.py --preset smoke --output results/smoke
python experiments/run_synthetic.py --preset full --output results/full
python scripts/build_paper.py
```

9 ICLR CHECKLIST

Anonymous submission. The paper and repository use anonymous author metadata.

Claims. The claims are limited to the synthetic benchmark and construction-level mechanism.

Code and data. All benchmark code is included. Data are generated procedurally.

Compute. Experiments run on CPU. The full preset uses 72 scenes and finishes in minutes on a standard laptop.

Limitations. The main limitation is lack of real neural-scene experiments.

REFERENCES

- Mildenhall, B., Srinivasan, P., Tancik, M., Barron, J., Ramamoorthi, R., and Ng, R. NeRF: Representing scenes as neural radiance fields for view synthesis. ECCV, 2020.
- Sitzmann, V., Zollhoefer, M., and Wetzstein, G. Scene representation networks: Continuous 3D-structure-aware neural scene representations. NeurIPS, 2019.
- Yu, A., Ye, V., Tancik, M., and Kanazawa, A. pixelNeRF: Neural radiance fields from one or few images. CVPR, 2021.
- Wang, Q., Wang, Z., Genova, K., Srinivasan, P., Zhou, H., Barron, J., Martin-Brualla, R., Snavely, N., and Funkhouser, T. IBRNet: Learning multi-view image-based rendering. CVPR, 2021.
- Niemeyer, M., Barron, J., Mildenhall, B., Sajjadi, M., Geiger, A., and Radwan, N. RegNeRF: Regularizing neural radiance fields for view synthesis from sparse inputs. CVPR, 2022.
- Wang, G., Chen, Z., Loy, C., and Liu, Z. SparseNeRF: Distilling depth ranking for few-shot novel view synthesis. ICCV, 2023.
- Kerbl, B., Kopanas, G., Leimkuehler, T., and Drettakis, G. 3D Gaussian Splatting for real-time radiance field rendering. SIGGRAPH, 2023.
- Guedon, A. and Lepetit, V. SuGaR: Surface-aligned Gaussian Splatting for efficient 3D mesh reconstruction. CVPR, 2024.
- Huang, B., Yu, Z., Chen, A., Geiger, A., and Gao, S. 2D Gaussian Splatting for geometrically accurate radiance fields. SIGGRAPH, 2024.
- Mescheder, L., Oechsle, M., Niemeyer, M., Nowozin, S., and Geiger, A. Occupancy Networks: Learning 3D reconstruction in function space. CVPR, 2019.
- Peng, S., Niemeyer, M., Mescheder, L., Pollefeys, M., and Geiger, A. Convolutional Occupancy Networks. ECCV, 2020.
- Cao, A. and de Charette, R. MonoScene: Monocular 3D semantic scene completion. CVPR, 2022.
- Li, Y., Yu, Z., Choy, C., Xiao, C., Alvarez, J., Fidler, S., Feng, C., and Anandkumar, A. VoxFormer: Sparse voxel transformer for camera-based 3D semantic scene completion. CVPR, 2023.
- Zheng, W. et al. OccWorld: Learning a 3D occupancy world model for autonomous driving. arXiv, 2024.
- Guzman-Rivera, A., Batra, D., and Kohli, P. Multiple choice learning: Learning to produce multiple structured outputs. NeurIPS, 2012.
- Bhattacharyya, A., Fritz, M., and Schiele, B. Best-of-many samples objective for generative models. CVPR, 2018.
- Cobbe, K., Kosaraju, V., Bavarian, M., Chen, M., Jun, H., Kaiser, L., Plappert, M., Tworek, J., Hilton, J., Nakano, R., Hesse, C., and Schulman, J. Training verifiers to solve math word problems. arXiv, 2021.

270 Lightman, H. et al. Let's verify step by step. ICLR, 2023.
271
272 Laurentini, A. The visual hull concept for silhouette-based image understanding. IEEE TPAMI,
273 1994.
274 Kutulakos, K. and Seitz, S. A theory of shape by space carving. IJCV, 2000.
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323